

Symbolic AI

Andre Freitas



Photo by Vasilyev Alexandr


Acknowledgements

- These slides were based on the slides of:
 - Peter A. Flach, Rule induction tutorial, IDA Spring School 2001.
 - Anoop & Hector, Inductive Logic Programming (for Dummies).
 - Gabor Melli, Scribe Notes on FOIL and Inverted Deduction.
 - CS 5751 Machine Learning, Chapter 10 Learning Sets of Rules.


This Lecture

- Introduction to Inductive Logic Programming
- FOIL


Linear Classifier




Decision Trees




Decision Trees



Rules



Clustering



ILP: Objective

Given a dataset:

- Positive examples ($E+$) and optionally negative examples ($E-$).
- Additional knowledge about the problem/application domain (Background Knowledge B).
- Set of constraints to make the learning process more efficient (C).

Goal of an ILP system is to find a set of hypothesis that:

- Explains (covers) the positive examples – Completeness.
- Are consistent with the negative examples – Consistency.

Generalisation & Specialisation

- **Generalising** a concept involves enlarging its extension in order to cover a given instance or subsume another concept.
- **Specialising** a concept involves restricting its extension in order to avoid covering a given instance or subsuming another concept.

First-order Representations

- **Propositional** representations:
 - datacase is *fixed-size vector of values*
 - features are those given in the dataset
- **First-order** representations:
 - datacase is *flexible-size, structured object*
 - sequence, set, graph
 - hierarchical: e.g. set of sequences
 - features need to be **selected** from potentially infinite set

Deductive Vs Inductive Reasoning

$T \cup B \rightarrow E$ (deduce)

parent(X,Y) :- mother(X,Y).
parent(X,Y) :- father(X,Y).

mother(mary,vinni).
mother(mary,andre).
father(carrey,vinni).
father(carrey,andre).

parent(mary,vinni).
parent(mary,andre).
parent(carrey,vinni).
parent(carrey,andre).

$E \cup B \rightarrow T$ (induce)

parent(mary,vinni).
parent(mary,andre).
parent(carrey,vinni).
parent(carrey,andre).

mother(mary,vinni).
mother(mary,andre).
father(carrey,vinni).
father(carrey,andre).

parent(X,Y) :- mother(X,Y).
parent(X,Y) :- father(X,Y).

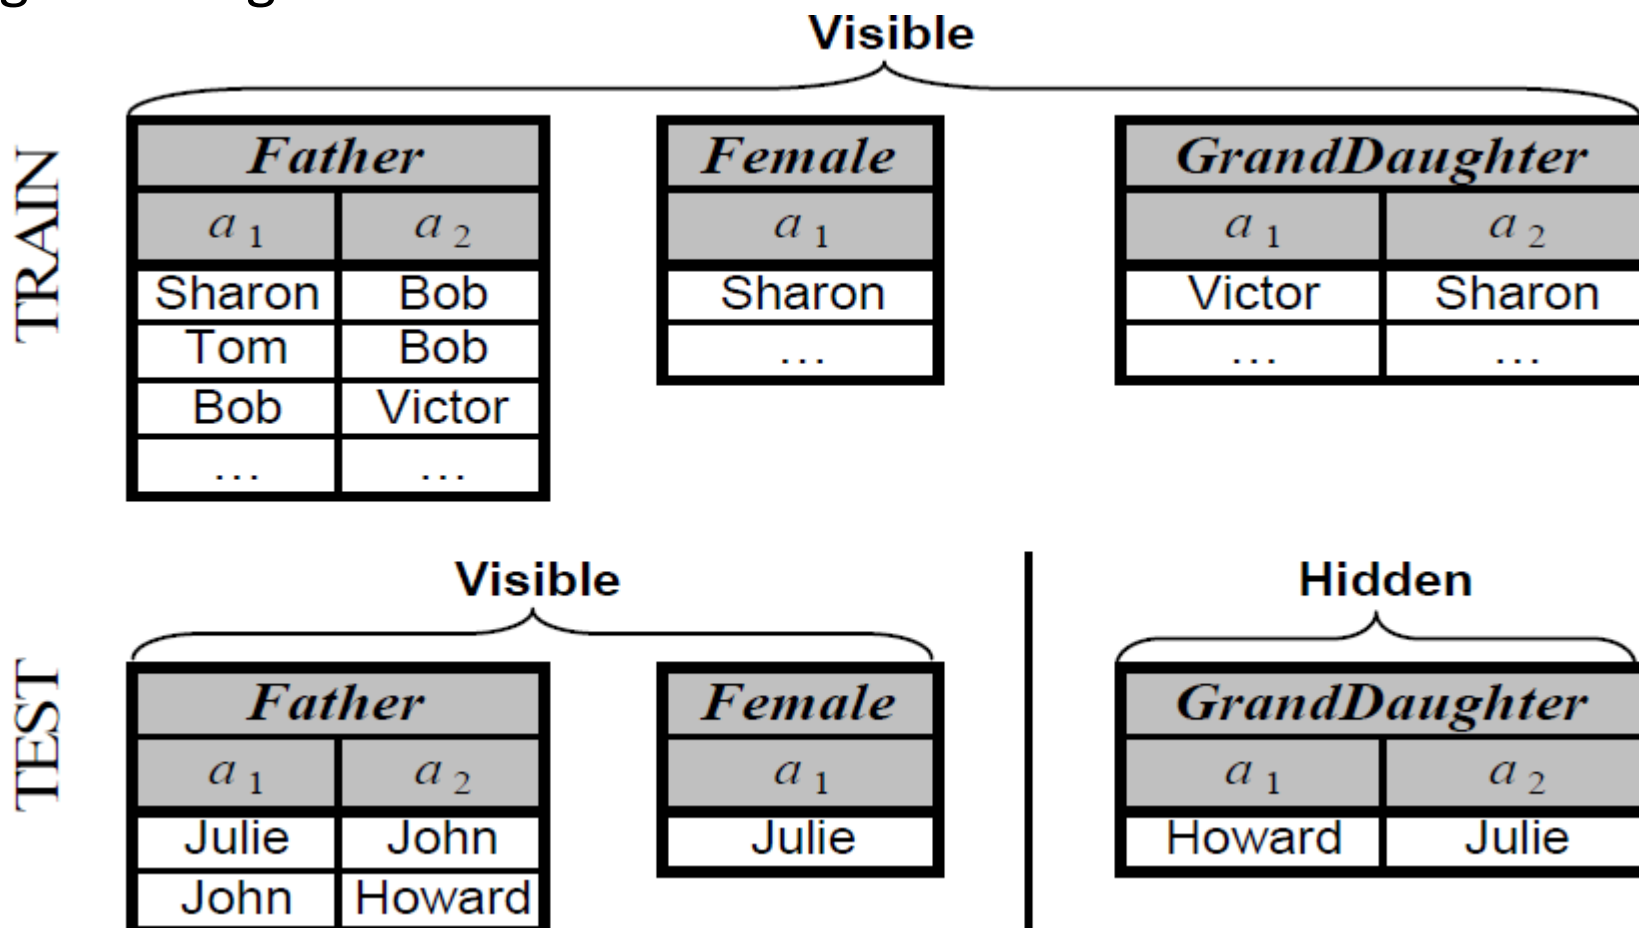
Relational Pattern

```
IF Customer(C1, Age1, Income1, TotSpent1, BigSpender1)
  AND MarriedTo(C1, C2)
  AND Customer(C2, Age2, Income2, TotSpent2, BigSpender2)
  AND Income2  $\geq$  10000
THEN BigSpender1 = Yes
```

```
big_spender(C1, Age1, Income1, TotSpent1)  $\leftarrow$ 
  married_to(C1, C2)  $\wedge$ 
  customer(C2, Age2, Income2, TotSpent2, BigSpender2)  $\wedge$ 
  Income2  $\geq$  10000
```

Example ILP Problem

Discover the rule that describes whether a person has a granddaughter



Propositional Learner with simple data transformation

- One of the first challenges that a propositional learner would encounter with this dataset is that the dataset is not structured as a set of fixed length-vectors of attribute-value pairs. This situation is typically resolved by JOINing the relations.

Predictors			Target
<i>Father</i>	<i>Child</i>	<i>Child is Fem.</i>	<i>Has Gdauh</i>
Bob	Sharon	TRUE	FALSE
Victor	Bob	FALSE	TRUE
...

Propositional Learner with simple data transformation

- A propositional learner would not locate a predictive model for this dataset.
- It would not be able to state that *Sharon* is *Victor's* granddaughter.
- At best it may discover that a child's gender has some influence on the likelihood that that child is a parent, or even a parent to a female child.

Predictors			Target
<i>Father</i>	<i>Child</i>	<i>Child is Fem.</i>	<i>Has Gdaugh</i>
Bob	Sharon	TRUE	FALSE
Victor	Bob	FALSE	TRUE
...

Propositional Learner with complex data transformation

- The algorithm cannot make the connection in one observation (*Bob* as a father) and another (*Bob* as child).
- A common way to enable a propositional learner to produce a predictive model on this data is to transform the data so that the required relations appear as attributes in the data.

- This transformation is sometimes referred to as 'flattening' the data.

Predictors					Target
<i>Father</i>	<i>Child</i>	<i>Child is Fem.</i>	<i>Child's Child</i>	<i>C's C is Fem.</i>	<i>Has Gdaugh</i>
Bob	Sharon	TRUE	NULL	NULL	FALSE
Victor	Bob	FALSE	Sharon	TRUE	TRUE
...

- Now the search for a rule is trivial. A decision tree would locate the pattern:

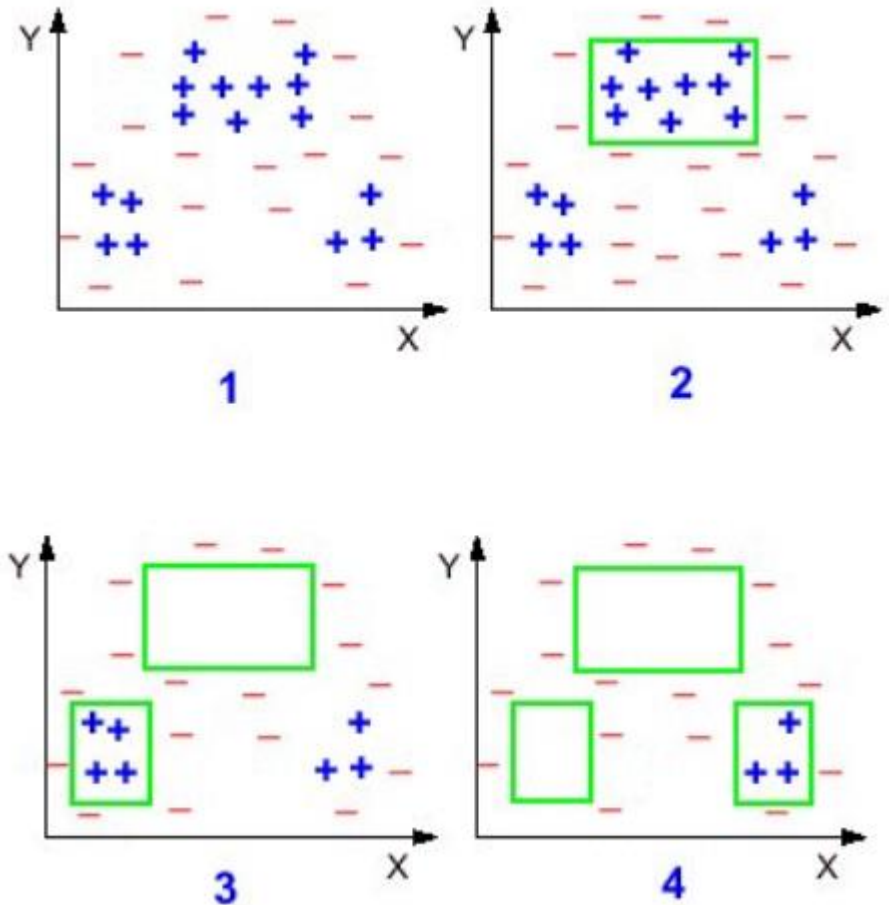
IF Child's Child is Female = TRUE
THEN HasGrandDaughter = TRUE.
ELSE HasGrandDaughter = FALSE

Propositional Sequential Covering

- A covering algorithm, in the context of propositional learning systems, is an algorithm that develops a cover for the set of positive examples.
 - that is, a set of hypotheses that account for all the positive examples but none of the negative examples.
- Sequential covering: it learns one rule at a time and repeat this process to gradually cover the full set of positive examples.

Iterate to Learn Multiple Rules

- Select seed from positive examples to build bottom clause.
- Get some rule “If $A \wedge B$ then P ”. Now throw away all positive examples that were covered by this rule
- Repeat until there are no more positive examples.



Propositional Sequential Covering

1. Start with an empty **Cover**
2. Use **Learn-One-Rule** to find the best hypothesis.
3. If the Just-Learnt-Rule satisfies the threshold then
 - Put Just-Learnt-Rule to the **Cover**.
 - Remove examples covered by Just-Learnt-Rule.
 - Go to step 2.
4. Sort the **Cover** according to its performance over examples.
5. Return: **Cover**.

Example

Id	Size	Colour	Shape	Weight	Expensive
1	Big	Red	Square	Heavy	Yes
2	Small	Blue	Triangle	Light	Yes
3	Small	Blue	Square	Light	No
4	Big	Green	Triangle	Heavy	No
5	Big	Blue	Square	Light	No
6	Big	Green	Square	Heavy	Yes
7	Small	Red	Triangle	Light	Yes

Expensive = Yes if:

Colour = Red.

(covers example 1,7)

Or (Colour = Green & Shape = Square).

(covers example 6)

Or (Colour = Blue & Shape = Triangle).

(covers example 2)

Complex

- A complex is a conjunction of attribute-value specifications. It forms the **condition** part in a rule, like "if **condition** then predict **class**".

Size=Big	Size=Small	Colour=Red
Colour=Green	Colour=Blue	Shape=Square
Shape=Triangle	Weight=Light	Weight=Heavy

- Specialising a complex is making a conjunction of the complex with one more attribute-value pair. For example:

Colour=Green & Shape=Square

(specialising Colour=Green or Shape=Square)

Colour=Blue & Weight=Heavy


(specialising Colour=Blue or Weight=Heavy)

Learn-One-Rule using Beam Search

1. Initialize a set of most general complexes.
2. Evaluate performances of those complexes over the example set.
 - Count how many positive and negative examples it covers.
 - Evaluate their performances.
3. Sort complexes according to their performances.
4. If the best complex satisfies some **threshold**, form the hypothesis and **return**.
5. Otherwise, pick k best performing complexes for the next generation.
6. Specializing all k complexes in the set to find new set of less general complexes.
7. Go to step 2.

The number k is the beam factor of the search, meaning the maximum number of complexes to be specialized.

Example



General to Specific Beam Search Example

- In the first step, 2 best complexes are found, namely $A=A1$ and $B=B2$.
- None of them satisfy the threshold, then the next level complexes are expanded and found 2 best complexes, eg. $A=A1 \ \& \ D=D1$ and $B=B2 \ \& \ A=A2$.
- The procedure keeps going until we find a complex that satisfies the threshold.

Entropy Evaluation Function

- The evaluation is based on the entropy of the set covered by that complex. Here is an example of a hypothesis covering 8 positive and 2 negative examples.

$$p1 = P(\text{positive}) = 8/(2+8) = 0.8;$$

$$p2 = P(\text{negative}) = 2/(2+8) = 0.2;$$

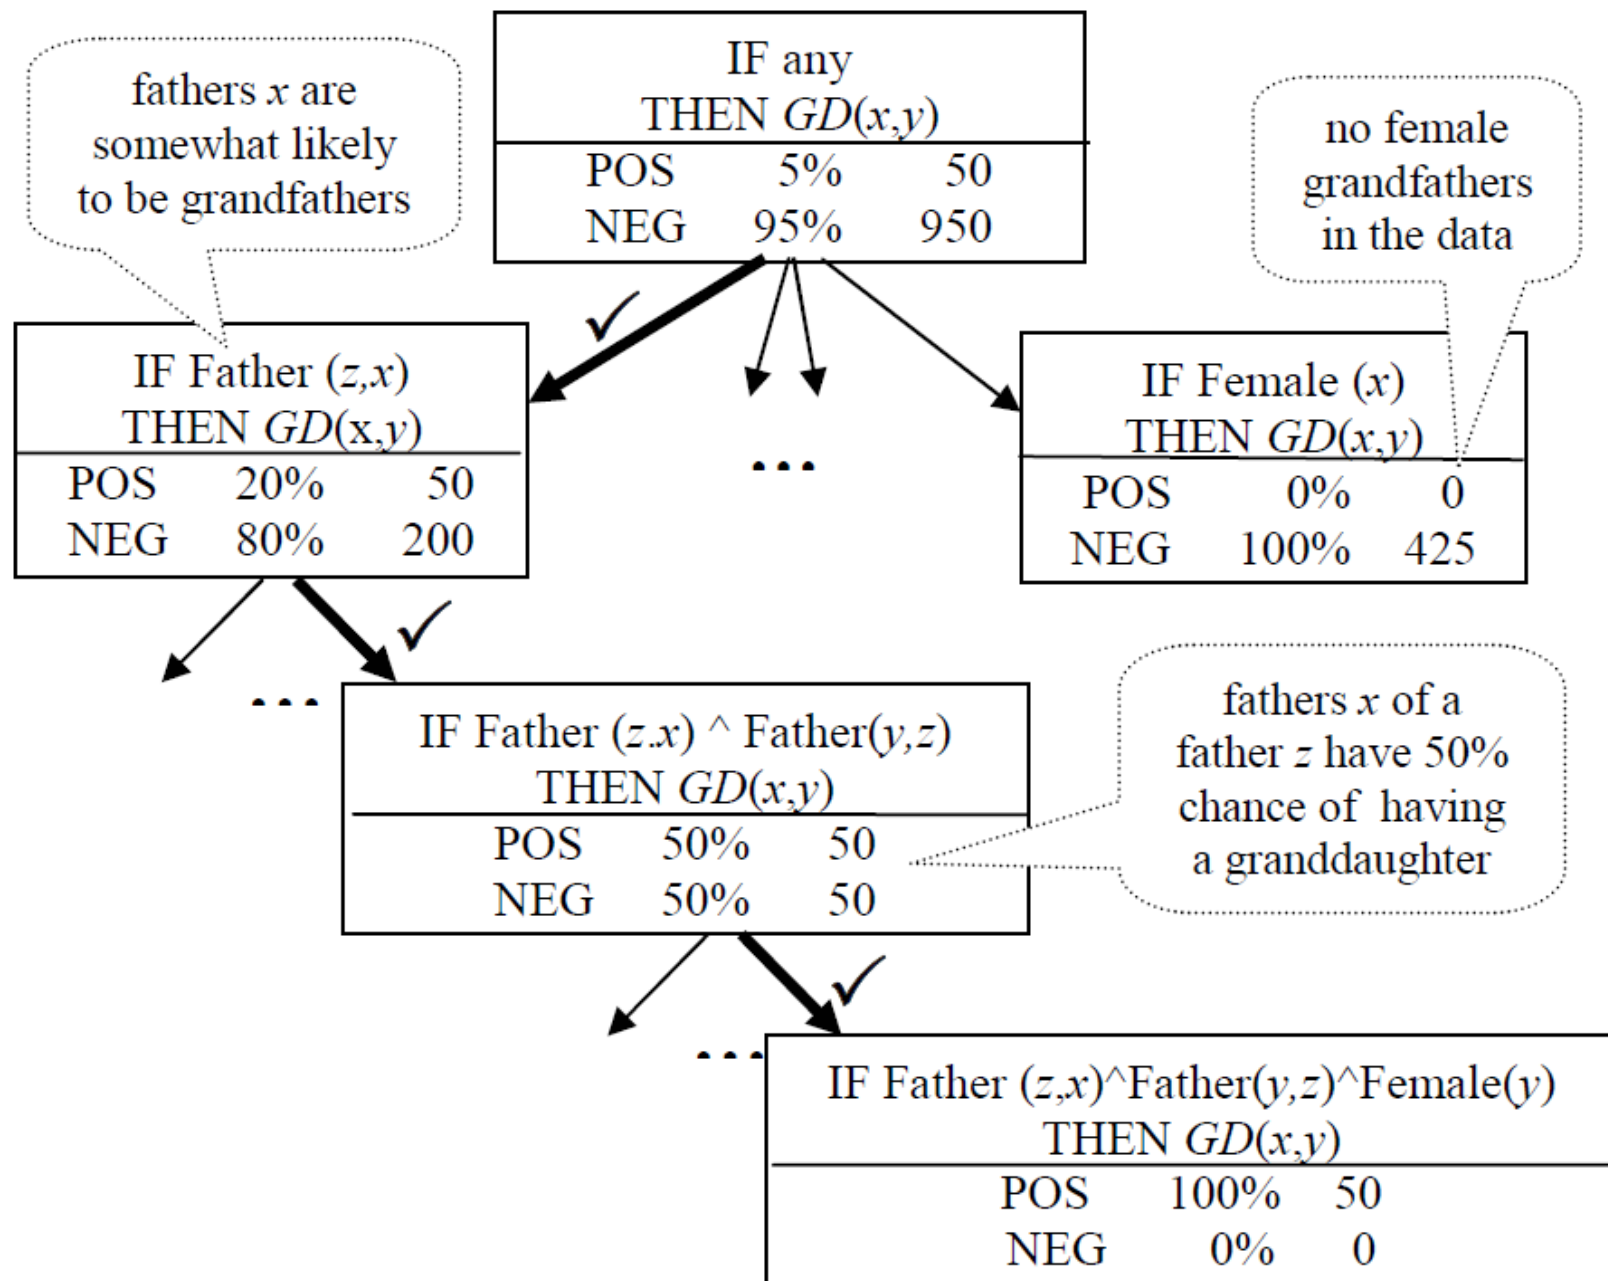
$$\text{Entropy} = - p1 * \log(p1) - p2 * \log(p2) = 0.72.$$

- In this function, the smaller the entropy is, the better the complex.
- In other words, the accuracy function can be defined as (1-Entropy).

The FOIL Algorithm

- The FOIL algorithm is a supervised learning algorithm that produces rules in first-order logic.
- The algorithm is a generalization of the SEQUENTIAL-COVERING and LEARN-ONE-RULE algorithms .
- The main modification is that search can also specialize on predicates with variables.
- The resulting rules differ from Horn clauses in two ways:
 - Negated symbols are allowed within the body.
 - FOIL's rules will not include function symbols.

Back to the Example



FOIL

FOIL(*Target_predicate*, *Predicates*, *Examples*)

Pos \leftarrow positive *Examples*

Neg \leftarrow negative *Examples*

while *Pos* do (*Learn a New Rule*)

NewRule \leftarrow most general rule possible

NegExamplesCovered \leftarrow *Neg*

while *NegExamplesCovered* do

 Add a new literal to specialize *NewRule*

 1. *Candidate_literals* \leftarrow generate candidates

 2. *Best_literal* $\leftarrow \operatorname{argmax}_{L \in \text{candidate_literal}} \text{FOIL_GAIN}(L, \text{NewRule})$

 3. Add *Best_literal* to *NewRule* preconditions

 4. *NegExamplesCovered* \leftarrow subset of *NegExamplesCovered* that satisfies *NewRule* preconditions

Learned_rules \leftarrow *Learned_rules* + *NewRule*

Pos \leftarrow *Pos* - {members of *Pos* covered by *NewRule*}

Return *Learned_rules*

The FOIL Algorithm

- The *outer loop adds new rules* to the output until no more positive examples are covered.
- The *inner loop searches for the next best rule* by incremental specialization.
- The outer loop corresponds to the SEQUENTIAL-CONVERGING algorithm, the inner to FIND-A-RULE

Specialising Rules in FOIL

Learning rule: $P(x_1, x_2, \dots, x_k) \leftarrow L_1 \dots L_n$

Candidate specializations add new literal of form:

- $Q(v_1, \dots, v_r)$, where at least one of the v_i in the created literal must already exist as a variable in the rule
- $Equal(x_j, x_k)$, where x_j and x_k are variables already present in the rule
- The negation of either of the above forms of literals

Information Gain in FOIL

$$FOIL_GAIN(L, R) \equiv t \left(\log_2 \frac{p_1}{p_1 + n_1} - \log_2 \frac{p_0}{p_0 + n_0} \right)$$

Where

- L is the candidate literal to add to rule R
- p_0 = number of positive bindings of R
- n_0 = number of negative bindings of R
- p_1 = number of positive bindings of $R+L$
- n_1 = number of negative bindings of $R+L$
- t is the number of positive bindings of R also covered by $R+L$

Note

- $-\log_2 \frac{p_0}{p_0 + n_0}$ is optimal number of bits to indicate the class of a positive binding covered by R

Applications

First Order Rule for Classifying Web Pages

From (Slattery, 1997)

course(A) \leftarrow
 has-word(A,instructor),
 NOT has-word(A,good),
 link-from(A,B)
 has-word(B,assignment),
 NOT link-from(B,C)

Train: 31/31, Test 31/34

Early diagnosis of rheumatic diseases

- Sample CN2 rule for an 8-class problem :

**IF Sex = male AND Age > 46 AND
Number_of_painful_joints > 3 AND
Skin_manifestations = psoriasis**

Application

A molecular compound is carcinogenic if:

- (1) it tests positive in the Salmonella assay; or
- (2) it tests positive for sex-linked recessive lethal mutation in Drosophila; or
- (3) it tests negative for chromosome aberration; or
- (4) it has a carbon in a six-membered aromatic ring with a partial charge of -0.13 ; or
- (5) it has a primary amine group and no secondary or tertiary amines; or
- (6) it has an aromatic (or resonant) hydrogen with partial charge ≥ 0.168 ; or
- (7) it has an hydroxy oxygen with a partial charge ≥ -0.616 and an aromatic (or resonant) hydrogen; or
- (8) it has a bromine; or
- (9) it has a tetrahedral carbon with a partial charge ≤ -0.144 and tests positive on Progol's mutagenicity rules.

**Final
Considerations**

Why ILP is not just Decision Trees.

- Language is First-Order Logic
 - Natural representation for multi-relational settings
 - Thus, a natural representation for *full* databases
- Not restricted to the classification task.
- So then, what is ILP?

Efficiency Issues

- Representational Aspects
- Search
- Evaluation
- Sharing computations
- Memory-wise scalability

Representational Aspects

- Example:
 - Student(string sname, string major, string minor)
 - Course(string cname, string prof, string cred)
 - Enrolled(string sname, string cname)
- In a natural join of these tables there is a one-to-one correspondence between join result and the Enrolled table.
- Data mining tasks on the Enrolled table are really propositional.

Representational Aspects

- Three settings for data mining:
 - Find patterns within individuals represented as tuples (single table, propositional)
 - eg. Which minor is chosen with what major
 - Find patterns within individuals represented as sets of tuples (each individual ‘induces’ a sub-database)
 - Multiple tables, restricted to some individual
 - eg. Student X taking course A, usually takes course B
 - Find patterns within the whole database
 - Multiple tables

Evaluation

- Evaluating a clause: get some measure of coverage
 - Match each example to the clause:
 - Run multiple logical queries.
 - Query optimization methods from DB community
 - Rel. Algebra operator reordering
 - BUT: queries for DB are set oriented (bottom-up), queries in PROLOG find a single solution (top-down).

Sharing Computations

- Materialization of features
- Propositionalization
- Pre-compute some statistics
 - Joint distribution over attributes of a table
 - Query selectivity
- Store proofs, reuse when evaluating new clauses

Summary

- Rules: easy to understand
 - Sequential covering algorithm
 - generate one rule at a time
 - general to specific - add antecedents
 - specific to general - delete antecedents
- First order logic and covering
 - how to connect variables
 - FOIL

Recommended Reading

QuickFOIL: Scalable Inductive Logic Programming

Qiang Zeng
University of
Wisconsin–Madison
qzeng@cs.wisc.edu

Jignesh M. Patel
University of
Wisconsin–Madison
jignesh@cs.wisc.edu

David Page
University of
Wisconsin–Madison
page@biostat.wisc.edu