

Virtualization

COMP 252 - Lecture 6

Antoniu Pop

antoniu.pop@manchester.ac.uk

16 February 2018

Previous Lecture: Virtualization Technologies

- ▶ **Aims of virtualization**

- ▶ Multiplex resources

Give the illusion that you own the resources.

- ▶ Isolation/abstraction

Software does not need to know the details of the hardware on which it runs.

- ▶ (avoid interference, safety, etc.)

- ▶ **Process vs. System Virtualization**


- ▶ **Process virtualization**

- ▶ JVM (“write once, run everywhere” model)
 - ▶ Dynamic Binary Translators (ISA: Rosetta, Mambo; OS&library calls: Wine)
 - ▶ Dynamic Binary Optimizers – program shepherding (Pin, Valgrind)


Today's Lecture – Learning Objectives

- ▶ To understand the implementation choices and details of System Virtualization
 - ▶ how virtualization works in modern architectures
 - ▶ what are the choices and characteristics of such implementations


Aims and Definitions



Hosted Virtualization



XEN Guest 0 Virtualization



Revision: OS Protection/Privilege

- ▶ OS handles physical resources
 - ▶ Privileged
- ▶ Application isolated from resources
 - ▶ Non-privileged

Application


Operating System

Hardware

Unvirtualized

Virtualization Protection/Privilege

- ▶ VMM handles physical resources
 - ▶ Privileged
- ▶ Guest OS isolated from resources
 - ▶ non- (or less-) privileged



Virtualized

VMM gets control on every guest OS access to physical resource



Guarded Physical Resources

- ▶ Timers
- ▶ CPU registers
 - ▶ Interrupt Enable
 - ▶ Page Table Base
- ▶ Device Control Registers
 - ▶ Programmed I/O?
 - ▶ Interrupt I/O?
 - ▶ DMA I/O?
- ▶ Interrupts (may be for different Guest?)
- ▶ Memory Mapping (page tables)

VMM Entry from Guest

- ▶ VMM designers are (a bit) lucky
 - ▶ Many Guest accesses to physical resources cause trap in non-privileged mode
 - ▶ So, running the OS in non-privileged mode suffices
- ▶ BUT some instructions behave differently (without trapping) in privileged and non-privileged mode (e.g. Intel “Store into Flags”)

Accessing Memory under Virtualization



What about TLBs?

Interfacing Guest OS and VMM

Three solutions today:

- ▶ Software (static)
- ▶ Software (dynamic)
- ▶ Hardware (dynamic)

ParaVirtualization

Modify Guest OS to be Virtualization-aware:

- ▶ call VMM for all privileged operations
- ▶ cooperate with VMM over shared page tables
- ▶ call VMM for input-output

Advantages? Disadvantages?

Detect and Fix Interfaces in VMM

► Detecting

- Write-protect Guest OS page tables
- Code-scan (Dynamic Binary Translation?) Guest OS for unsafe instructions – plant traps

► Fixing

- Use write-error trap to detect guest page-table writes
- Provide “shadow page tables” for hardware TLBs
- Use “illegal instruction” and “trap” traps

Detect and Fix Interfaces in Hardware

- ▶ Requirement
 - ▶ VMM runs more-privileged than Guest OS
- ▶ Hardware provides Application/OS and VMM modes
- ▶ When Virtualization is active, all OS accesses to physical resources trap to VMM

Advantages? Disadvantages?