

COMP26120

Academic Session: 2018-19

Lab Exercise 10: Testing the small world hypothesis

Duration: 3 lab sessions. **WARNING: You will find it difficult to complete this exercise if you do not work for all three weeks.**

For this lab exercise you should do all your work in your COMP26120/ex10 directory.

For this assignment, you can only get full credit if you do not use code from outside sources. You can get partial credit by using properly attributed code from outside sources (see marking scheme).

Learning Objectives

By the end of this lab you should be able to:

- Explain the all-pairs shortest path problem in graphs and describe two algorithms for computing it and their properties
- Implement a priority queue data structure and explain the complexity of its main operations
- Implement Dijkstra's algorithm, and use it to test the small world hypothesis on social network data.
- Implement a heuristic path finding algorithm using local information, apply it to the same data and then compare the results.
- Use complexity arguments to show which of these two is more efficient for a particular set of data.

Introduction

The “small world” hypothesis

We are all parts of social networks. These are the networks of people we interact with socially: acquaintances, friends, colleagues, teachers, family, etc. We can think of these networks as graphs. Each individual is a node, and is connected to the people that individual personally knows.

One property these social networks are thought to have is the “small-world property” that any pair of nodes in this very large network is connected by a short path. It is often said that any two humans are connected by a path involving no more than five other people, so the small-world property is also referred to as “[six degrees of separation](#)”.

The six degrees of separation notion comes from a set of [experiments](#) carried out by Stanley Milgram in the late 1960s, which attempted to investigate chains of acquaintances in the United States. In the experiment, Milgram sent several packages to 160 random people living in Omaha, Nebraska, asking them to give the package to a friend or acquaintance whom they thought would bring the package closer to a set final individual, a stockbroker living in Boston, Massachusetts, 1450 miles away. Many were not delivered, but those that were passed through on average about six hands from sender to receiver. The experiment itself has since been criticised, but the idea that social networks possess this special small-world property has persisted.

If Milgram's experiments are to be believed, social networks graphs must possess two properties.

1. Between any two nodes there must be at least one short path (the small-world property).
2. Assuming that not all paths are short, there must be some information available locally at each node which suggests a good next step, i.e one that is part of a short path.

This lab comes in three parts. Parts 1 and 2 test the first property: are the paths short in particular social networks. Part 3 tests the second property: is there a simple way to find the next node in a short path.

The social network data

The social network data is found in `/opt/info/courses/COMP26120/problems/ex10/`. These are the networks of Facebook friends from two US universities, Caltech and University of Oklahoma. These are represented as graphs in the same format as the graphs from lab exercise 9, and are in the files *Caltech.gx* and *Oklahoma.gx*. The Caltech data contains 769 nodes and 33,312 edges (e.g. approximately 43 friends in the same university each). The Oklahoma data contains 17,425 nodes and 1,785,056 edges (approximately 102 friends each). In both cases, only the friends from within the same university are included. The data was collected on a particular date in September 2005.

It should be noted that a Facebook friends network may be very different from a real social network. Unfortunately, it is very difficult to get data on real social networks, because this information is not recorded in any one place.

Description

Part 1: Background Investigation

This part is worth 5 marks in total

Copy into your ex10 directory the file *LabReport.tex* from the usual *COMP26120/problems/ex10* directory. (Don't copy the data files. They are about 50M and may crunch your quota.) Write in the first section of the *LabReport.tex* file your statement of the small-world hypothesis and how you are going to test it (without describing the algorithms).

We need a way to find the shortest path between the nodes of a graph, and we are interested in the shortest paths between all of the nodes. This is called the all-pairs shortest path problem. There are (at least) two algorithms to solve this: Dijkstra's algorithm and Floyd's algorithm (also called Floyd-Warshall's algorithm).

For these graphs, Dijkstra's algorithm is more efficient. We would like you to learn about the complexity of these two algorithms and find out why Dijkstra's algorithm is more efficient for these graphs. Write the complexities of the two algorithms in the *LabReport.tex* document and the argument showing that Dijkstra's algorithm is more efficient for these graphs. Note that you need to find

*Tejas Chandrasekar*²

the complexity of Dijkstra's algorithm for the all-pairs shortest path problem, not the single-source shortest path problem.

Part 2: Finding the the shortest paths

This part is worth 15 marks in total

The graphs should be represented in the same format as in Lab exercise 9, so you can use the same code to read in the data files.

Part 2a: Implementation of the shortest-path algorithm

You will need to learn about the Dijkstra's algorithm and implement it. Dijkstra's algorithm requires a priority-first search. It contains a graph traversal algorithm which is similar to those you may have implemented in lab exercise 9, except **you don't pull from your search list the most recently added (depth-first search) or the one which has been longest in the list (breadth-first search)**, but the one which is **closest in distance to the starting node**. Thus, you need to implement a **priority queue**.

Updated: The remainder of Part 2a has been updated since the lab was first published. The social network given in this lab is not a weighted graph. You should process length of each edge in your algorithm as one equal unit. It has been commented by some students that in an unweighted graph there is no need for a *priority* queue. Therefore, there are two options for the next part:

1. *Implement a priority queue.* You can implement a priority queue inefficiently, by using a linear search to find the best item each time you pull from the queue. You will get more points if you implement this in one of the efficient ways. For instance, it can be implemented as a heap. The relevant section in the textbook, is section 2.4, and particularly 2.4.3. Finding shortest paths in graphs is discussed in Chapter 7 of the textbook. (If you use code from outside sources for this, you can only get partial credit for this part. See marking scheme for details.)
2. *Implement a queue.* You can use a simple queue implementation **if** you also provide a justification for why this is correct for unweighted graphs in your *LabReport.tex*. The explanation should argue in terms of the workings of Dijkstra's algorithm. You will be asked to explain this justification during marking but you **must** have written out a justification in your report.

After implementing the (priority) queue, you can complete the implementation of Dijkstra's algorithm and put it all in a file called *part1.c*. You should use and extend *graph.h* and *graph_functions.c* from exercise 9.

Part 2b: Test for the small-world property

Update: This section has been updated to be more open and useful.

Use your algorithm to measure properties of the paths in the two social networks to test the hypothesis. You should collect enough statistics to be able to answer the hypothesis. It is not good enough to just collect statistics without knowing what they mean. Record the statistics and your conclusions from them in the *LabReport.tex* file. Are these small-world networks?

As well as these statistics, collect and record execution times. Given the number of edges and vertices in each graph, do your execution times agree with the complexity characteristics found in Part 1?

Part 3: Approximate path finding and heuristics

This part is worth 10 marks in total

Once an all-pairs shortest path algorithm is run, it is possible to give each node a look-up table of the next node in the shortest path to any target node (a routing table). However, sometimes it is not feasible to run an algorithm which explores the entire graph to generate this table.

Certainly the participants in Milgram's experiment could only know about their friends and acquaintances, and had to choose which of those to give it to based on some quantity used to determine which of those was most likely to be the next step on the shortest path, or at least a short path. We will refer to such a quantity as a heuristic. The heuristic they might have used may have been geographical (give it to someone who has some connection to somewhere close to Boston), or social (give it to someone who knows a lot of people).

Using heuristics to find approximations to shortest paths may be useful when the graph is too big to implement Dijkstra or Floyd, or when it is changing too rapidly for there to be up-to-date information about the entire network at any node. In this part, you will investigate whether such heuristics find short paths in these networks.

We will call this approach "heuristic path following". It works like this. You choose some heuristic. To find a path to a target, each node chooses the next node from its outlist based on this heuristic. For instance, for the Facebook data, the heuristic could be, choose the node among its friends with the most friends (the friendship relation is mutual, so out-degree is the same as in-degree in these graphs). Using this heuristic, the algorithm works as follows, to find the path from SOURCE to TARGET:

```
1 CURRENT ← SOURCE;
2 while (TARGET not in CURRENT.OUTLIST) and (CURRENT.OUTLIST not empty)
3     add CURRENT to PATH
4     Let MAXOUT be the unvisited node in CURRENT.OUTLIST with largest out-
       degree
5     CURRENT ← MAXOUT;
6 endwhile
7 add CURRENT to PATH
```

Figure 1: heuristic path following

Here CURRENT is the current node, and PATH stores the entire path. This approach is clearly not guaranteed to find the shortest paths (why not?).

Your task is to implement this algorithm and test to what extent it finds shortest paths. Put the code in a file called *part2.c*.

Design an experiment to test how accurate the heuristic method is and use it to compare this approach with that of Part 2 (actual shortest paths) on the two data sets. Write the results and conclusions in the *LabReport.tex* document.

Submission and Marking Scheme

You should submit *LabReport.pdf*, *graph.h*, *graph_functions.c*, *part1.c* and *part2.c*.

This is a large lab and we will not be able to mark everybody in one session. Please have patience and attend early labs otherwise we may run out of time for marking overall.

Part 1. Is the small-world hypothesis clearly stated along with plans on how to test it?	
The hypothesis is clearly stated and one or more well-thought out experiments have been described which will help establish whether the hypothesis holds. The student can clearly explain both the hypothesis and tests.	(2)
As above but the tests are not very well described in the report.	(1.5)
The hypothesis is stated and some tests exist but the student cannot clearly explain them or they are flawed e.g. they do not properly test the hypothesis	(1)
No attempt has been made	(0)

Part 1. Are the complexities of the two algorithms stated and explained?	
Both complexities for the all-paths shortest paths algorithms are given, there is a written explanation of why Dijkstra's algorithm is better for these graphs (referring directly to qualities of the given graphs), and the student can explain their reasoning.	(3)
The complexities are given and there is an explanation but it is a little confused or not quite precise enough.	(2)
The answers are both good but the student cannot adequately explain those answers	(1)
Either the complexities or explanation has a significant mistake in it but a reasonable attempt has been made	(1)
No attempt has been made	(0)

Part 2a. Is Dijkstra's algorithm completely and correctly implemented (ignoring for now the use of a priority queue or otherwise)?	
The code is complete (all necessary functionality is there) and looks correct. The student may be asked to demonstrate correctness on the small.gx file from ex9. It is expected that the student has performed tests on this graph or others of their own to establish correctness.	(3)
There is a minor mistake in the algorithm.	(2)
There is a major mistake in the algorithm or it is incomplete.	(1)
No attempt has been made	(0)

Has a priority queue been implemented or an explanation given for why it is not necessary for unweighted graphs	
A good implementation of a priority queue has been provided utilising an efficient method (e.g. a heap). The student can explain how it works.	(3)
A correct but inefficient implementation of a priority queue has been provided (e.g. as a linear search). The student can explain how it works.	(2)
A priority queue implementation is present but the student cannot explain how it works.	(1)
A clear explanation of why a priority queue is not needed in this case is provided. This directly refers to how Dijkstra's algorithm works. A queue is provided instead and its functionality can be explained.	(3)
An explanation of why a priority queue is not needed is provided but it is vague. A queue is provided instead and its functionality can be explained.	(2)
No priority queue has been implemented and no justification has been given	(0)
No attempt has been made	(0)

Part 2a. Can the student explain how Dijkstra's algorithm works?	
The student has excellent understanding. They can explain the general operation of the algorithm. They can explain the role of the priority queue in the weighted graph case. They can explain why the algorithm will always find the shortest path.	(3)
The student has good understanding. They can explain the general operation of the algorithm.	(2)
The student has a high-level understanding of how the algorithm works but struggles to explain the details.	(1)
No attempt has been made	(0)

Part 2b. Has the hypothesis been tested and the results and conclusions clearly reported?	
A detailed report of the tests carried out and the results found is given in the Lab Report. This should include raw data from the tests and analysis of those results before reaching a conclusion. The tests should be reproducible by running the submitted code and the analysis should be sound.	(3)
Some tests have been carried out and used to reach a conclusion. There should be some evidence of thinking about the data to reach the conclusion (e.g. analysis).	(2)
Some test data is present but it is just stated without any analysis.	(1)
No attempt has been made	(0)

Part 2b. Have execution times been recorded and used to explore the complexity of the algorithm?	
Detailed execution times have been recorded. This will be for more than one graph and times may have been recorded for subparts of the algorithm. The properties of the different graphs (e.g. number of vertices and edges, connectness, sparsity) will be used to discuss whether the theoretical complexities hold in practice.	(3)
Execution times are recorded and there is a discussion of what they might mean with respect to complexities but the analysis is shallow.	(2)
Execution times are merely recorded and stated.	(1)
No attempt has been made	(0)

Part 3. The heuristic path following algorithm has been implemented correctly and completely.	
The algorithm is complete and correct. It will always terminate (e.g. eventually reach the target or return an error if it does not exist) and will not visit the same node more than once. The given heuristic does not need to be used, e.g. the student can try out their own, but the chosen heuristic should be correctly applied.	(3)
As above but there is a minor mistake	(2)
As above but there is a major mistake	(1)
No attempt has been made	(0)

Part 3. The student understands their code and the nature of heuristic search in general.	
The student's understanding is excellent. The student can explain their code. The student can explain why this approach is not guaranteed to find the shortest path. The student can reflect on why a greedy approach works in Dijkstra's algorithm but not here.	(3)
The student's understanding is good. They can explain their code and attempt answers to the other questions.	(2)
The student attempts an explanation but the topic needs revision	(1)
No attempt has been made	(0)

Has a test been designed and carried out to investigate the effectiveness of the heuristic method?	
A sound (it makes sense) and significant (there is enough data) test has been designed and implemented. The raw results have been collected and reported. These results have been analysed and a conclusion has been reached. The test should make a direct comparison with the solution provided in Part 2.	(4)
As above but there are some minor issues	(3)
Relevant data has been collected but there is not sufficient analysis to draw a conclusion	(2)
Some data has been collected but it is not relevant to answering the question	(1)
No attempt has been made	(0)